

AUTOMATED RATING REVIEW SYSTEM

Model Documentation

NAÏVE BAYES

1. Model Used

The model used throughout this project is the Multinomial Naive Bayes (MNB) classifier. It is a probabilistic machine learning model based on Bayes' Theorem and is particularly well-suited for text classification tasks where features represent word counts or TF-IDF scores.

2. Why Multinomial Naive Bayes?

- Efficient for Text Data: Works very well with high-dimensional sparse features produced by TF-IDF vectorization.
- Computationally Fast: Requires less training time compared to models like SVM or RandomForest.
- Good Baseline Model: Provides reliable performance for sentiment and rating prediction tasks.
- Handles Imbalanced Data: Performs reasonably well even when some star ratings are less frequent.

3. Working Principle

Based on Bayes' Theorem:

$$P(\text{Class}|\text{Words}) = [P(\text{Words}|\text{Class}) * P(\text{Class})] / P(\text{Words})$$

It assumes that the features (words) are independent of each other given the class. For each review, the model calculates the probability of it belonging to each rating class (1★ to 5★) and assigns the class with the highest probability.

4. Why Not Other Models?

- Random Forest / Gradient Boosting: Very powerful for structured/numeric datasets but not efficient for sparse text data.
- Support Vector Machine (SVM): Can achieve high accuracy but computationally expensive for large datasets and requires tuning.

5. Strengths of Naive Bayes in This Project

- Fast to train and predict.
- Works effectively with TF-IDF features.
- Requires minimal parameter tuning.
- Provides interpretable results (class probabilities).

6. Limitations

- Independence assumption (words treated as unrelated) is not always realistic.
- May struggle when features are highly correlated.
- Usually outperformed by deep learning models for very complex tasks.

7. Hyperparameter Tuning

In this project, we performed **hyperparameter tuning** for the **Naïve Bayes classifier** using **GridSearchCV**.

- **Hyperparameter tuned:** alpha (Laplace smoothing)
- **Purpose:** Controls smoothing to handle zero probabilities for unseen words in text data. Lower values make the model less smoothed, which can improve predictions for frequent words while reducing the effect on rare/unseen words.
- **Method used:** GridSearchCV from scikit-learn
 - Tested multiple values: [0.1, 0.5, 1.0, 1.5, 2.0]
 - Evaluated using cross-validation to select the best performing value.
- **Best value found:** alpha = 0.1
- **Impact:** Using the tuned alpha slightly improved the model's performance, especially for minority classes in both balanced and imbalanced datasets.

7. Conclusion

In this project, Multinomial Naive Bayes was chosen because it is simple yet powerful for text-based problems, scalable for large datasets, efficient in computation, and well-suited for review rating prediction tasks. Thus, all model training, evaluation, and inference steps were carried out using Naive Bayes as the core model.

Codes :

1. Balance the dataset

```
# Balance the dataset to have equal number of reviews per rating
balanced_df = balance_dataset(review, target_col="Rating", n_samples=5000)
```

2. Train-Test Split & TF-IDF Vectorization

```
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer

X = balanced_df['Review_text']
y = balanced_df['Rating']

# Stratified train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, stratify=y, random_state=42
)

# TF-IDF Vectorization
tfidf = TfidfVectorizer(max_features=5000, ngram_range=(1,2), stop_words='english')
X_train_tfidf = tfidf.fit_transform(X_train)
X_test_tfidf = tfidf.transform(X_test)
```

3. Train Naïve Bayes Model

```
from sklearn.naive_bayes import MultinomialNB

# Initialize and train the model
nb_model = MultinomialNB()
nb_model.fit(X_train_tfidf, y_train)
```

4. Evaluate the Model and Hyper parameter tuning

```
# Hyperparameter tuning
param_grid = {'alpha': [0.1, 0.5, 1.0, 1.5, 2.0]}
grid = GridSearchCV(MultinomialNB(), param_grid, cv=5, scoring='accuracy')
grid.fit(X_train_tfidf, y_train)

best_nb_model = grid.best_estimator_
print("Best alpha:", grid.best_params_)

# Evaluate tuned model
y_pred = best_nb_model.predict(X_test_tfidf)
print("Naive Bayes Accuracy (tuned):", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))
```